

Retail CapstoneCP

June 17, 2023

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from datetime import timedelta
from pandas import ExcelWriter
```

```
[2]: df = pd.read_excel("Online Retail.xlsx")
df.head()
```

```
[2]: InvoiceNo StockCode Description Quantity \
0 536365 85123A WHITE HANGING HEART T-LIGHT HOLDER 6
1 536365 71053 WHITE METAL LANTERN 6
2 536365 84406B CREAM CUPID HEARTS COAT HANGER 8
3 536365 84029G KNITTED UNION FLAG HOT WATER BOTTLE 6
4 536365 84029E RED WOOLLY HOTTIE WHITE HEART. 6
```

```
InvoiceDate UnitPrice CustomerID Country
0 2010-12-01 08:26:00 2.55 17850.0 United Kingdom
1 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
2 2010-12-01 08:26:00 2.75 17850.0 United Kingdom
3 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
4 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
```

```
[3]: # Check shape of data
df.shape
```

```
[3]: (541909, 8)
```

```
[4]: # Check feature details of data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	InvoiceNo	541909 non-null	object
1	StockCode	541909 non-null	object
2	Description	540455 non-null	object
3	Quantity	541909 non-null	int64
4	InvoiceDate	541909 non-null	datetime64[ns]
5	UnitPrice	541909 non-null	float64
6	CustomerID	406829 non-null	float64
7	Country	541909 non-null	object

dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 33.1+ MB

```
[5]: # Check missing values in data
df.isnull().sum()
```

```
[5]: InvoiceNo      0
      StockCode    0
      Description  1454
      Quantity     0
      InvoiceDate   0
      UnitPrice    0
      CustomerID  135080
      Country      0
      dtype: int64
```

```
[6]: # Calculating the Missing Values % contribution in DF
df_null = round(df.isnull().sum()/len(df)*100,2)
df_null
```

```
[6]: InvoiceNo      0.00
      StockCode    0.00
      Description  0.27
      Quantity     0.00
      InvoiceDate   0.00
      UnitPrice    0.00
      CustomerID  24.93
      Country      0.00
      dtype: float64
```

As we can see two columns in data have missing values. * Description - 0.27% (1454 nos.) *
CustomerID - 24.93% (135080)

CustomerID is important feature of our analysis since our analysis is centered around Customers only so we can not impute null values **CustomerID** with mean/ median/ mode in this case. We will check possibility to fill null values in **CustomerID** column by looking up for **InvoiceNo** of the row having null **CustomerID** in other rows where **CustomerID** is present. If there are still any null values in **CustomerID** after this process then we will drop complete row having missing

CustomerID.

We can drop **Description** feature from our data since it is not going to contribute in our model.

```
[7]: invoice_null_custid = set(df[df['CustomerID'].isnull()][ 'InvoiceNo'])
df[df['InvoiceNo'].isin(invoice_null_custid) & (~df['CustomerID'].isnull())]
```

```
[7]: Empty DataFrame
Columns: [InvoiceNo, StockCode, Description, Quantity, InvoiceDate, UnitPrice,
CustomerID, Country]
Index: []
```

We could not find any value to impute null values in **CustomerID** column since all entries for a particular **InvoiceNo** have missing **CustomerID** if that particular **InvoiceNo** has null **CustomerID** in even one entry. So we will drop all rows having null values in **CustomerID**.

```
[8]: df = df.drop('Description', axis=1)
df = df.dropna()
df.shape
```

```
[8]: (406829, 7)
```

- **(b) Remove duplicate data records:** Since our data is transactional data and it has duplicate entries for InvoiceNo and CustomerID, we will drop only those rows which are completely duplicated, not on the basis of any one particular column such as InvoiceNo or CustomerID etc.

```
[9]: df = df.drop_duplicates()
df.shape
```

```
[9]: (401602, 7)
```

- **(c) Perform descriptive analysis on the given data:**

```
[10]: # CustomerID is 'float64', changing the datatype of CustomerId to string as
      ↪ Customer ID as numerical data does not make sense

df['CustomerID'] = df['CustomerID'].astype(str)
```

```
[11]: df.describe(datetime_is_numeric=True)
```

```
[11]:
```

	Quantity	InvoiceDate	UnitPrice
count	401602.000000	401602	401602.000000
mean	12.182579	2011-07-10 12:08:08.129743104	3.474064
min	-80995.000000	2010-12-01 08:26:00	0.000000
25%	2.000000	2011-04-06 15:02:00	1.250000
50%	5.000000	2011-07-29 15:40:00	1.950000
75%	12.000000	2011-10-20 11:58:00	3.750000

max	80995.000000	2011-12-09 12:50:00	38970.000000
std	250.283248	NaN	69.764209

- **Quantity:** Average quantity of each product in transaction is 12.18. Also note that minimum value in **Quantity** column is negative. This implies that some customers had returned the product during our period of analysis.
- **InvoiceDate:** Our data has transaction between 01-12-2010 to 09-12-2011
- **UnitPrice:** Average price of each product in transactions is 3.47

```
[12]: df.describe(include=['O'])
```

```
[12]:
```

	InvoiceNo	StockCode	CustomerID	Country
count	401602	401602	401602	401602
unique	22190	3684	4372	37
top	576339	85123A	17841.0	United Kingdom
freq	542	2065	7812	356726

- **InvoiceNo:** Total entries in preprocessed data are 4,01,602 but transactions are 22,190. Most number of entries (count of unique products) are in Invoice No. '576339' and is 542 nos.
- **StockCode:** There are total 3684 unique products in our data and product with stock code '85123A' appears most frequently (2065 times) in our data.
- **CustomerID:** There are 4372 unique customers in our final preprocessed data. Customer with ID '17841' appears most frequently in data (7812 times)
- **Country:** Company has customers across 37 countries. Most entries are from United Kingdom in our dataset (356726)

0.0.1 (B) Data Transformation

(2) Perform Cohort Analysis * (a) Create month cohort of customers and analyze active customers in each cohort:

```
[13]: # Convert to InvoiceDate to Year-Month format
df['month_year'] = df['InvoiceDate'].dt.to_period('M')
df['month_year'].nunique()
```

```
[13]: 13
```

```
[14]: month_cohort = df.groupby('month_year')['CustomerID'].nunique()
month_cohort
```

```
[14]: month_year
2010-12    948
2011-01    783
2011-02    798
2011-03   1020
2011-04    899
2011-05   1079
```

```

2011-06    1051
2011-07     993
2011-08     980
2011-09    1302
2011-10    1425
2011-11    1711
2011-12     686
Freq: M, Name: CustomerID, dtype: int64

```

```


[15]: plt.figure(figsize=(10,5))
      sns.barplot(y = month_cohort.index, x = month_cohort.values);
      plt.xlabel("Count of customers")
      plt.title("No. of active customers in each month")

```

```

[15]: Text(0.5, 1.0, 'No. of active customers in each month')

```



- (b) Analyze the retention rate of customers:

```

[16]: month_cohort - month_cohort.shift(1)

```

```

[16]: month_year
2010-12    NaN
2011-01   -165.0
2011-02    15.0
2011-03   222.0
2011-04  -121.0
2011-05   180.0
2011-06  -28.0

```

```

2011-07    -58.0
2011-08    -13.0
2011-09    322.0
2011-10    123.0
2011-11    286.0
2011-12   -1025.0
Freq: M, Name: CustomerID, dtype: float64

```


```
[17]: retention_rate = round(month_cohort.pct_change(periods=1)*100,2)
      retention_rate
```

```

[17]: month_year
2010-12      NaN
2011-01   -17.41
2011-02     1.92
2011-03    27.82
2011-04   -11.86
2011-05    20.02
2011-06    -2.59
2011-07    -5.52
2011-08    -1.31
2011-09    32.86
2011-10     9.45
2011-11    20.07
2011-12   -59.91
Freq: M, Name: CustomerID, dtype: float64

```

```
[18]: plt.figure(figsize=(10,5))
      sns.barplot(y = retention_rate.index, x = retention_rate.values);
      plt.xlabel("Retention (in %)")
      plt.title("Month-wise customer retention rate");
```



0.1 Week 2:

Monetary analysis:

```
[20]: df['amount'] = df['Quantity']*df['UnitPrice']
df.head()
```

```
[20]: InvoiceNo StockCode Quantity InvoiceDate UnitPrice CustomerID \
0 536365 85123A 6 2010-12-01 08:26:00 2.55 17850.0
1 536365 71053 6 2010-12-01 08:26:00 3.39 17850.0
2 536365 84406B 8 2010-12-01 08:26:00 2.75 17850.0
3 536365 84029G 6 2010-12-01 08:26:00 3.39 17850.0
4 536365 84029E 6 2010-12-01 08:26:00 3.39 17850.0
```

```
Country month_year amount
0 United Kingdom 2010-12 15.30
1 United Kingdom 2010-12 20.34
2 United Kingdom 2010-12 22.00
3 United Kingdom 2010-12 20.34
4 United Kingdom 2010-12 20.34
```

```
[21]: df_monetary = df.groupby('CustomerID').sum()['amount'].reset_index()
df_monetary
```

```
[21]: CustomerID amount
0 12346.0 0.00
1 12347.0 4310.00
```

2	12348.0	1797.24
3	12349.0	1757.55
4	12350.0	334.40
...
4367	18280.0	180.60
4368	18281.0	80.82
4369	18282.0	176.60
4370	18283.0	2045.53
4371	18287.0	1837.28

[4372 rows x 2 columns]

Frequency Analysis:

```
[22]: df_frequency = df.groupby('CustomerID').nunique()['InvoiceNo'].reset_index()
# df_frequency = df.drop_duplicates('InvoiceNo').groupby('CustomerID').
# count()['InvoiceNo'].reset_index()
df_frequency
```

```
[22]:      CustomerID  InvoiceNo
0         12346.0          2
1         12347.0          7
2         12348.0          4
3         12349.0          1
4         12350.0          1
...         ...         ...
4367        18280.0          1
4368        18281.0          1
4369        18282.0          3
4370        18283.0         16
4371        18287.0          3
```

[4372 rows x 2 columns]

Recency Analysis:

```
[23]: # We will fix reference date for calculating recency as last transaction day in
# data + 1 day
ref_day = max(df['InvoiceDate']) + timedelta(days=1)
df['days_to_last_order'] = (ref_day - df['InvoiceDate']).dt.days
df.head()
```

```
[23]:      InvoiceNo  StockCode  Quantity      InvoiceDate  UnitPrice  CustomerID  \
0      536365      85123A          6 2010-12-01 08:26:00          2.55      17850.0
1      536365       71053          6 2010-12-01 08:26:00          3.39      17850.0
2      536365      84406B          8 2010-12-01 08:26:00          2.75      17850.0
3      536365      84029G          6 2010-12-01 08:26:00          3.39      17850.0
4      536365      84029E          6 2010-12-01 08:26:00          3.39      17850.0
```


	Country	month_year	amount	days_to_last_order
0	United Kingdom	2010-12	15.30	374
1	United Kingdom	2010-12	20.34	374
2	United Kingdom	2010-12	22.00	374
3	United Kingdom	2010-12	20.34	374
4	United Kingdom	2010-12	20.34	374

```
[24]: df_recency = df.groupby('CustomerID')['days_to_last_order'].min().reset_index()
df_recency
```

```
[24]:
```

	CustomerID	days_to_last_order
0	12346.0	326
1	12347.0	2
2	12348.0	75
3	12349.0	19
4	12350.0	310
...
4367	18280.0	278
4368	18281.0	181
4369	18282.0	8
4370	18283.0	4
4371	18287.0	43

[4372 rows x 2 columns]

Calculate RFM metrics:

```
[25]: df_rf = pd.merge(df_recency, df_frequency, on='CustomerID', how='inner')
df_rfm = pd.merge(df_rf, df_monetary, on='CustomerID', how='inner')
df_rfm.columns = ['CustomerID', 'Recency', 'Frequency', 'Monetary']
df_rfm.head()
```


```
[25]:
```

	CustomerID	Recency	Frequency	Monetary
0	12346.0	326	2	0.00
1	12347.0	2	7	4310.00
2	12348.0	75	4	1797.24
3	12349.0	19	1	1757.55
4	12350.0	310	1	334.40

Build RFM Segments:


```
[26]: df_rfm['recency_labels'] = pd.cut(df_rfm['Recency'], bins=5,
                                         labels=['newest', 'newer', 'medium', 'old',
                                         ↪ 'older', 'oldest'])
df_rfm['recency_labels'].value_counts().plot(kind='barh');
df_rfm['recency_labels'].value_counts()
```

```
[26]: newest      2734
      newer       588
      medium     416
      older      353
      oldest     281
      Name: recency_labels, dtype: int64
```




```
[27]: df_rfm['frequency_labels'] = pd.cut(df_rfm['Frequency'], bins=5,
      ↪ labels=['lowest', 'lower', 'medium', 'higher', 'highest'])
      df_rfm['frequency_labels'].value_counts().plot(kind='barh');
      df_rfm['frequency_labels'].value_counts()
```

```
[27]: lowest      4348
      lower       18
      medium       3
      highest      2
      higher       1
      Name: frequency_labels, dtype: int64
```



```
[28]: df_rfm['monetary_labels'] = pd.cut(df_rfm['Monetary'], bins=5,  
    ↪ labels=['smallest', 'smaller', 'medium', 'larger', 'largest'])  
df_rfm['monetary_labels'].value_counts().plot(kind='barh');  
df_rfm['monetary_labels'].value_counts()
```

```
[28]: smallest    4357  
smaller        9  
medium         3  
largest        2  
larger         1  
Name: monetary_labels, dtype: int64
```



```
[29]: df_rfm['rfm_segment'] =
    ↪ df_rfm[['recency_labels', 'frequency_labels', 'monetary_labels']].agg('-',
    ↪ join, axis=1)
df_rfm.head()
```

```
[29]: CustomerID  Recency  Frequency  Monetary  recency_labels  frequency_labels \
0    12346.0      326         2         0.00      oldest      lowest
1    12347.0         2         7    4310.00      newest      lowest
2    12348.0        75         4    1797.24      newest      lowest
3    12349.0        19         1    1757.55      newest      lowest
4    12350.0       310         1     334.40      oldest      lowest
```

```
monetary_labels      rfm_segment
0      smallest  oldest-lowest-smallest
1      smallest  newest-lowest-smallest
2      smallest  newest-lowest-smallest
3      smallest  newest-lowest-smallest
4      smallest  oldest-lowest-smallest
```

RFM Score:

```
[30]: recency_dict = {'newest': 5, 'newer':4, 'medium': 3, 'older':2, 'oldest':1}
frequency_dict = {'lowest':1, 'lower':2, 'medium': 3, 'higher':4, 'highest':5}
monetary_dict = {'smallest':1, 'smaller':2, 'medium': 3, 'larger':4, 'largest':
    ↪ 5}
```


```
df_rfm['rfm_score'] = df_rfm['recency_labels'].map(recency_dict).astype(int)+  
↳df_rfm['frequency_labels'].map(frequency_dict).astype(int) +  
↳df_rfm['monetary_labels'].map(monetary_dict).astype(int)  
df_rfm.head(10)
```

```
[30]: CustomerID Recency Frequency Monetary recency_labels frequency_labels \  
0 12346.0 326 2 0.00 oldest lowest  
1 12347.0 2 7 4310.00 newest lowest  
2 12348.0 75 4 1797.24 newest lowest  
3 12349.0 19 1 1757.55 newest lowest  
4 12350.0 310 1 334.40 oldest lowest  
5 12352.0 36 11 1545.41 newest lowest  
6 12353.0 204 1 89.00 medium lowest  
7 12354.0 232 1 1079.40 older lowest  
8 12355.0 214 1 459.40 medium lowest  
9 12356.0 23 3 2811.43 newest lowest
```


```
monetary_labels rfm_segment rfm_score  
0 smallest oldest-lowest-smallest 3  
1 smallest newest-lowest-smallest 7  
2 smallest newest-lowest-smallest 7  
3 smallest newest-lowest-smallest 7  
4 smallest oldest-lowest-smallest 3  
5 smallest newest-lowest-smallest 7  
6 smallest medium-lowest-smallest 5  
7 smallest older-lowest-smallest 4  
8 smallest medium-lowest-smallest 5  
9 smallest newest-lowest-smallest 7
```

Analyze RFM Segment and Score:

```
[32]: df_rfm['rfm_segment'].value_counts().plot(kind='barh', figsize=(10, 5));
```



```
[33]: df_rfm['rfm_score'].value_counts().plot(kind='barh', figsize=(10, 5));
```



0.2 Week 3

0.2.1 Data Modeling:

1. Create clusters using k-means clustering algorithm.

a. Prepare the data for the algorithm. If the data is asymmetrically distributed, manage the skewness with appropriate transformation. Standardize the data.

```
[34]: print(df_rfm.shape)
df_rfm.head()
```

(4372, 9)

```
[34]:
```

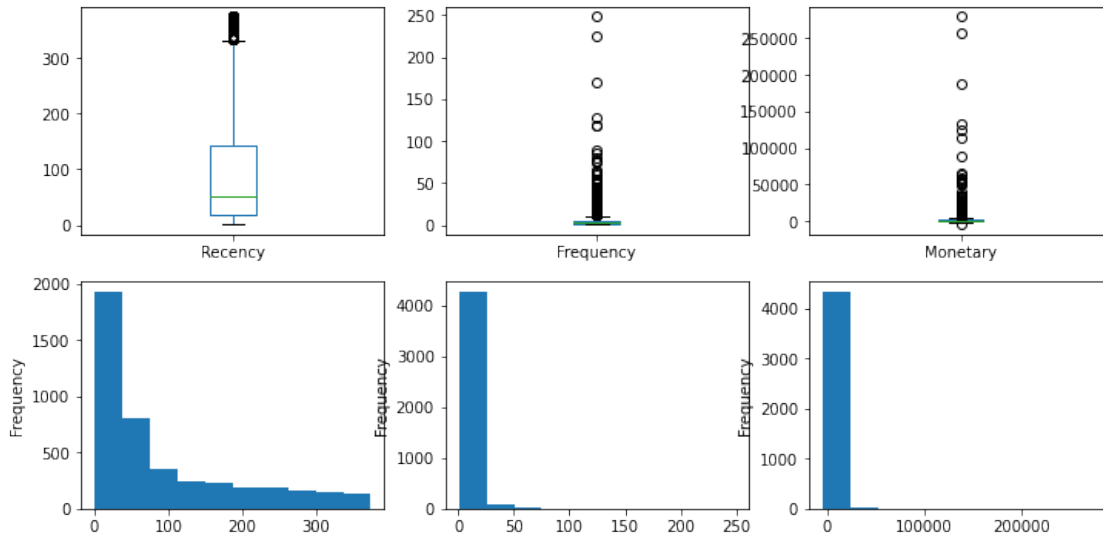
	CustomerID	Recency	Frequency	Monetary	recency_labels	frequency_labels	\
0	12346.0	326	2	0.00	oldest	lowest	
1	12347.0	2	7	4310.00	newest	lowest	
2	12348.0	75	4	1797.24	newest	lowest	
3	12349.0	19	1	1757.55	newest	lowest	
4	12350.0	310	1	334.40	oldest	lowest	

	monetary_labels	rfm_segment	rfm_score
0	smallest	oldest-lowest-smallest	3
1	smallest	newest-lowest-smallest	7
2	smallest	newest-lowest-smallest	7

```
3         smallest  newest-lowest-smallest      7
4         smallest  oldest-lowest-smallest     3
```

```
[36]: plt.figure(figsize=(12,6))

for i, feature in enumerate(['Recency', 'Frequency', 'Monetary']):
    plt.subplot(2,3,i+1)
    df_rfm[feature].plot(kind='box')
    plt.subplot(2,3,i+1+3)
    df_rfm[feature].plot(kind='hist')
```



0.3 Week 3

0.3.1 Data Modeling:

1. Create clusters using k-means clustering algorithm.

a. Prepare the data for the algorithm. If the data is asymmetrically distributed, manage the skewness with appropriate transformation. Standardize the data.

Outliers: Frequency and Monetary features in above data seem to have lot of outliers. Lets drop them.

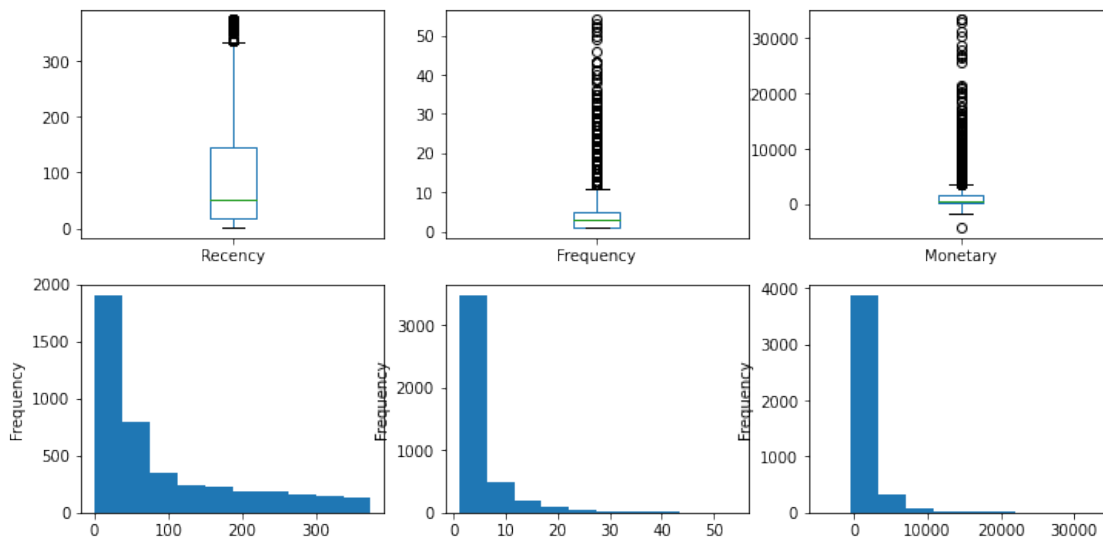
```
[37]: df_rfm = df_rfm[(df_rfm['Frequency']<60) & (df_rfm['Monetary']<40000)]
df_rfm.shape
```

```
[37]: (4346, 9)
```

26 Customers removed as outlier from out data.

```
[39]: plt.figure(figsize=(12,6))

for i, feature in enumerate(['Recency', 'Frequency', 'Monetary']):
    plt.subplot(2,3,i+1)
    df_rfm[feature].plot(kind='box')
    plt.subplot(2,3,i+1+3)
    df_rfm[feature].plot(kind='hist')
```



Log Transformation: Now since all three features have right skewed data therefore we will use log transformation of these features in our model.

```
[40]: df_rfm_log_trans = pd.DataFrame()
df_rfm_log_trans['Recency'] = np.log(df_rfm['Recency'])
df_rfm_log_trans['Frequency'] = np.log(df_rfm['Frequency'])
df_rfm_log_trans['Monetary'] = np.log(df_rfm['Monetary']-df_rfm['Monetary'].
    ↪min()+1)
```

Standard Scalar Transformation: It is extremely important to rescale the features so that they have a comparable scale.

```
[41]: scaler = StandardScaler()

df_rfm_scaled = scaler.fit_transform(df_rfm_log_trans[['Recency', 'Frequency',
    ↪'Monetary']])
df_rfm_scaled

df_rfm_scaled = pd.DataFrame(df_rfm_scaled)
df_rfm_scaled.columns = ['Recency', 'Frequency', 'Monetary']
df_rfm_scaled.head()
```



```
[41]:      Recency  Frequency  Monetary
0  1.402988  -0.388507  -0.770922
1 -2.100874   0.967301   1.485132
2  0.392218   0.361655   0.364190
3 -0.552268  -1.138669   0.342970
4  1.368370  -1.138669  -0.527416
```

b. Build K-Means Clustering Model and Decide the optimum number of clusters to be formed.

```
[42]: # k-means with some arbitrary k
kmeans = KMeans(n_clusters=3, max_iter=50)
kmeans.fit(df_rfm_scaled)
```

```
[42]: KMeans(max_iter=50, n_clusters=3)
```


```
[43]: kmeans.labels_
```

```
[43]: array([2, 1, 0, ..., 0, 1, 0], dtype=int32)
```

```
[45]: # Finding the Optimal Number of Clusters with the help of Elbow Curve/ SSD
ssd = []
range_n_clusters = [2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12]
for num_clusters in range_n_clusters:
    kmeans = KMeans(n_clusters=num_clusters, max_iter=100)
    kmeans.fit(df_rfm_scaled)

    ssd.append(kmeans.inertia_)

# plot the SSDs for each n_clusters
plt.plot(range_n_clusters,ssd);
```



```
[46]: # Creating dataframe for exporting to create visualization in tableau later
df_inertia = pd.DataFrame(list(zip(range_n_clusters, ssd)),
    columns=['clusters', 'intertia'])
df_inertia
```

```
[46]:
```

	clusters	intertia
0	2	7113.079519
1	3	5343.136928
2	4	4480.972122
3	5	3730.838474
4	6	3044.921651
5	7	2598.297835
6	8	2299.191284
7	9	2044.597189
8	10	1852.964018
9	11	1700.376798
10	12	1575.623062

```
[47]: # Finding the Optimal Number of Clusters with the help of Silhouette Analysis
range_n_clusters = [2, 3, 4, 5, 6, 7, 8, 9, 10]

for num_clusters in range_n_clusters:
    kmeans = KMeans(n_clusters=num_clusters, max_iter=50)
    kmeans.fit(df_rfm_scaled)

    cluster_labels = kmeans.labels_
```

```

silhouette_avg = silhouette_score(df_rfm_scaled, cluster_labels)
print("For n_clusters={0}, the silhouette score is {1}".
      ↪format(num_clusters, silhouette_avg))

```

```

For n_clusters=2, the silhouette score is 0.44132753537785846
For n_clusters=3, the silhouette score is 0.38135409490825667
For n_clusters=4, the silhouette score is 0.3623606426972478
For n_clusters=5, the silhouette score is 0.36479947836332627
For n_clusters=6, the silhouette score is 0.34419116171743463
For n_clusters=7, the silhouette score is 0.34288879825359414
For n_clusters=8, the silhouette score is 0.3354507509941627
For n_clusters=9, the silhouette score is 0.3463017984588029
For n_clusters=10, the silhouette score is 0.3560796733393901

```

```

[48]: # Final model with k=3
kmeans = KMeans(n_clusters=3, max_iter=50)
kmeans.fit(df_rfm_scaled)

```

```

[48]: KMeans(max_iter=50, n_clusters=3)

```

c. Analyze these clusters and comment on the results.

```

[49]: # assign the label
df_rfm['Cluster_Id'] = kmeans.labels_
df_rfm.head()

```

```

[49]:  CustomerID  Recency  Frequency  Monetary  recency_labels  frequency_labels  \
0    12346.0      326         2         0.00         oldest         lowest
1    12347.0         2         7    4310.00         newest         lowest
2    12348.0        75         4    1797.24         newest         lowest
3    12349.0        19         1    1757.55         newest         lowest
4    12350.0       310         1     334.40         oldest         lowest

```

```


      monetary_labels      rfm_segment  rfm_score  Cluster_Id
0      smallest  oldest-lowest-smallest         3         1
1      smallest  newest-lowest-smallest         7         2
2      smallest  newest-lowest-smallest         7         0
3      smallest  newest-lowest-smallest         7         1
4      smallest  oldest-lowest-smallest         3         1

```


```

[50]: # Box plot to visualize Cluster Id vs Monetary
sns.boxplot(x='Cluster_Id', y='Monetary', data=df_rfm);


```



```
[51]: # Box plot to visualize Cluster Id vs Frequency
sns.boxplot(x='Cluster_Id', y='Frequency', data=df_rfm);
```



```
[52]: # Box plot to visualize Cluster Id vs Recency
sns.boxplot(x='Cluster_Id', y='Recency', data=df_rfm);
```



0.3.2 Inference:

As we can observe from above boxplots that our model has nicely created 3 segments of customer with the interpretation as below: * Customers with Cluster Id 0 are less frequent buyers with low monetary expenditure and also they have not purchased anything in recent time and hence least important for business. * Customers with Cluster Id 1 are the customers having Recency, Frequency and Monetary score in the medium range. * Customers with Cluster Id 2 are the most frequent buyers, spending high amount and recently placing orders so they are the most important customers from business point of view.

0.4 Week 4:

Data Reporting: 1. Create a dashboard in tableau by choosing appropriate chart types and metrics useful for the business. The dashboard must entail the following:

- Country-wise analysis to demonstrate average spend. Use a bar chart to show the monthly
- Bar graph of top 15 products which are mostly ordered by the users to show the number of
- Bar graph to show the count of orders vs. hours throughout the day
- Plot the distribution of RFM values using histogram and frequency charts
- Plot error (cost) vs. number of clusters selected
- Visualize to compare the RFM values of the clusters using heatmap

```
[ ]: # Writing dataframe to excel file for creating visualization in tableau
writer = pd.ExcelWriter('C:\\Users\\mgupt\\mgpython\\Capstone Project\\Retail ->PGP\\output_data.xlsx', engine='xlsxwriter')

df.to_excel(writer, sheet_name='master_data', index=False)
df_rfm.to_excel(writer, sheet_name='rfm_data', index=False)
df_inertia.to_excel(writer, sheet_name='inertia', index=False)
writer.save()

[ ]: product_desc = pd.read_excel("Online Retail.xlsx")
product_desc = product_desc[['StockCode', 'Description']]
product_desc = product_desc.drop_duplicates()
product_desc.to_csv('product_desc.csv', index=False)

[ ]:
```